Beauty And The Burst

Remote Identification of Encrypted Video Streams

Presented by: Aviel Zecharia & Eli Cohen

Agenda

- Background
- MPEG-DASH Standard
- Attack Overview
- Deep Neural Networks
- Adversarial Models
  - On-Path Attack
  - Off-Path Attack

Authors

Roei Schuster

- Tel Aviv University
- Cornell Tech

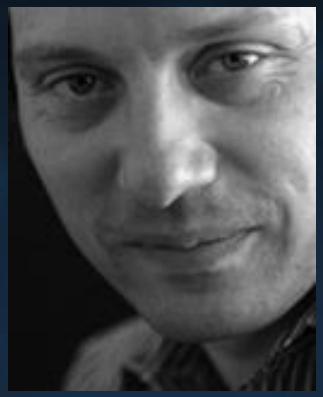

Vitaly Shmatikov

- Cornell Tech

Eran Tromer

- Tel Aviv University
- Columbia University

Why video traffic is so interesting ?

Background

- Targeted marketing purposes
- Market characterization efforts
- Not everybody wants to volunteer this information about their habits
  - Video traffic is encrypted
  - HTTPS has been in wide deployment

What still can be learned ?

Initial buffering & on/off bursts

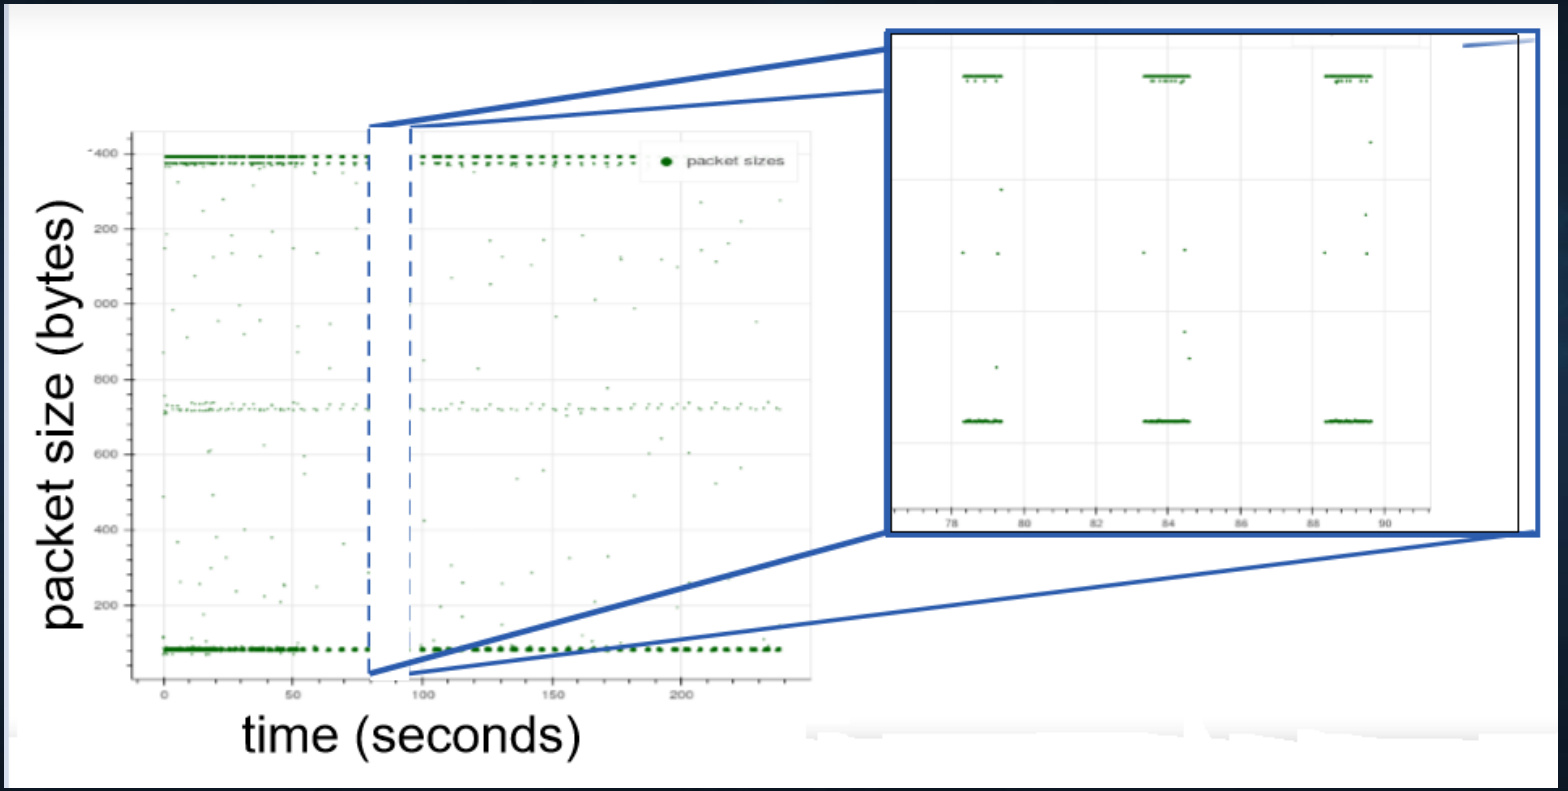

Where the bursts come from ?

MPEG-DASH standard

- Widely adopted by major streaming providers
  - Netflix, YouTube etc.
- Variable bitrate (VBR)

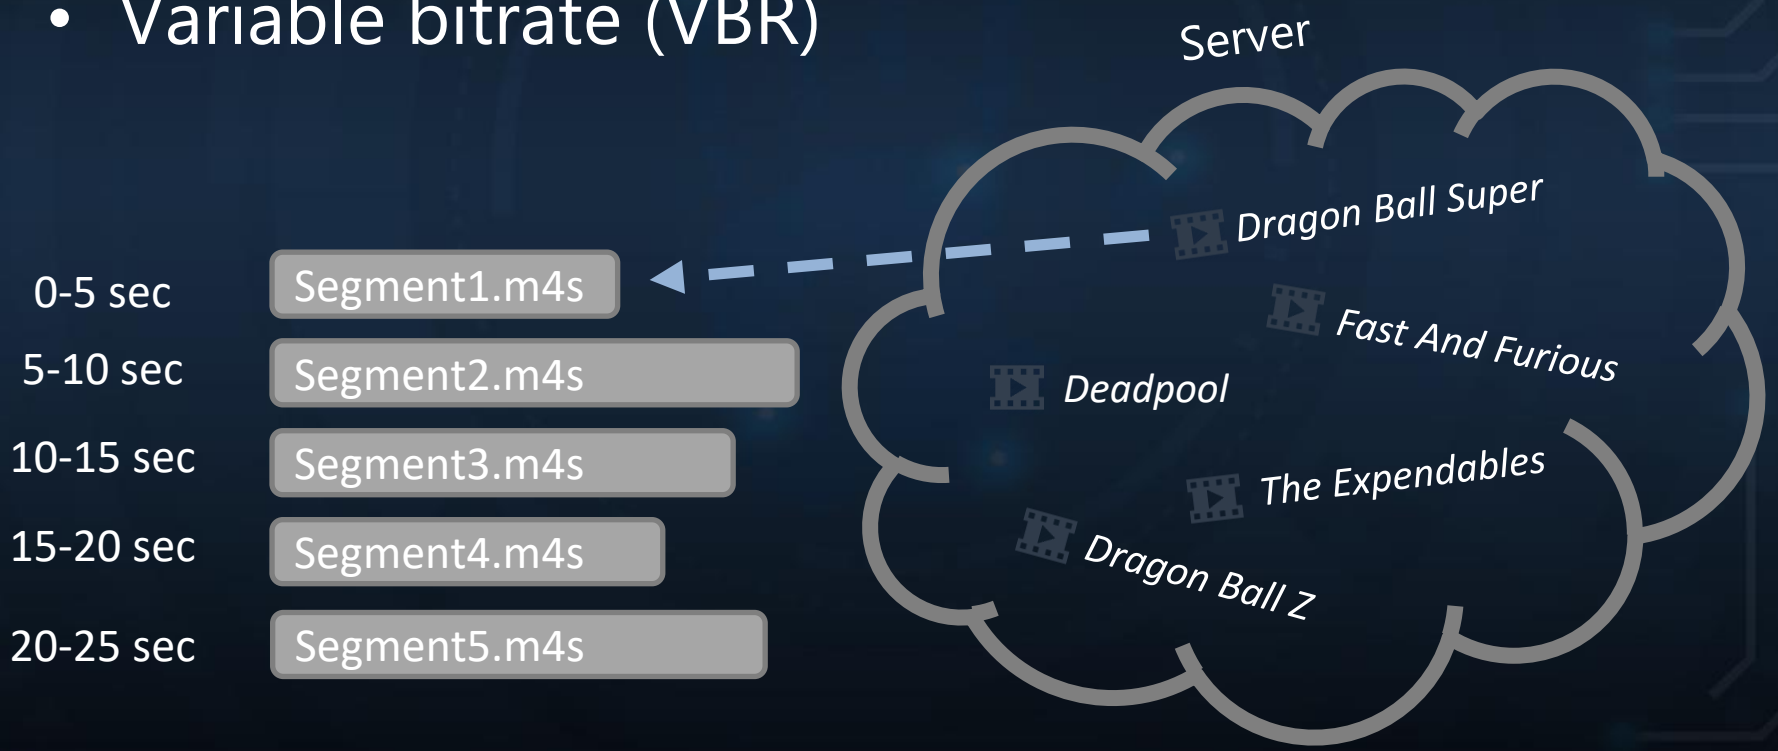

MPEG-DASH standard

- Adaptive bitrate streaming over HTTP

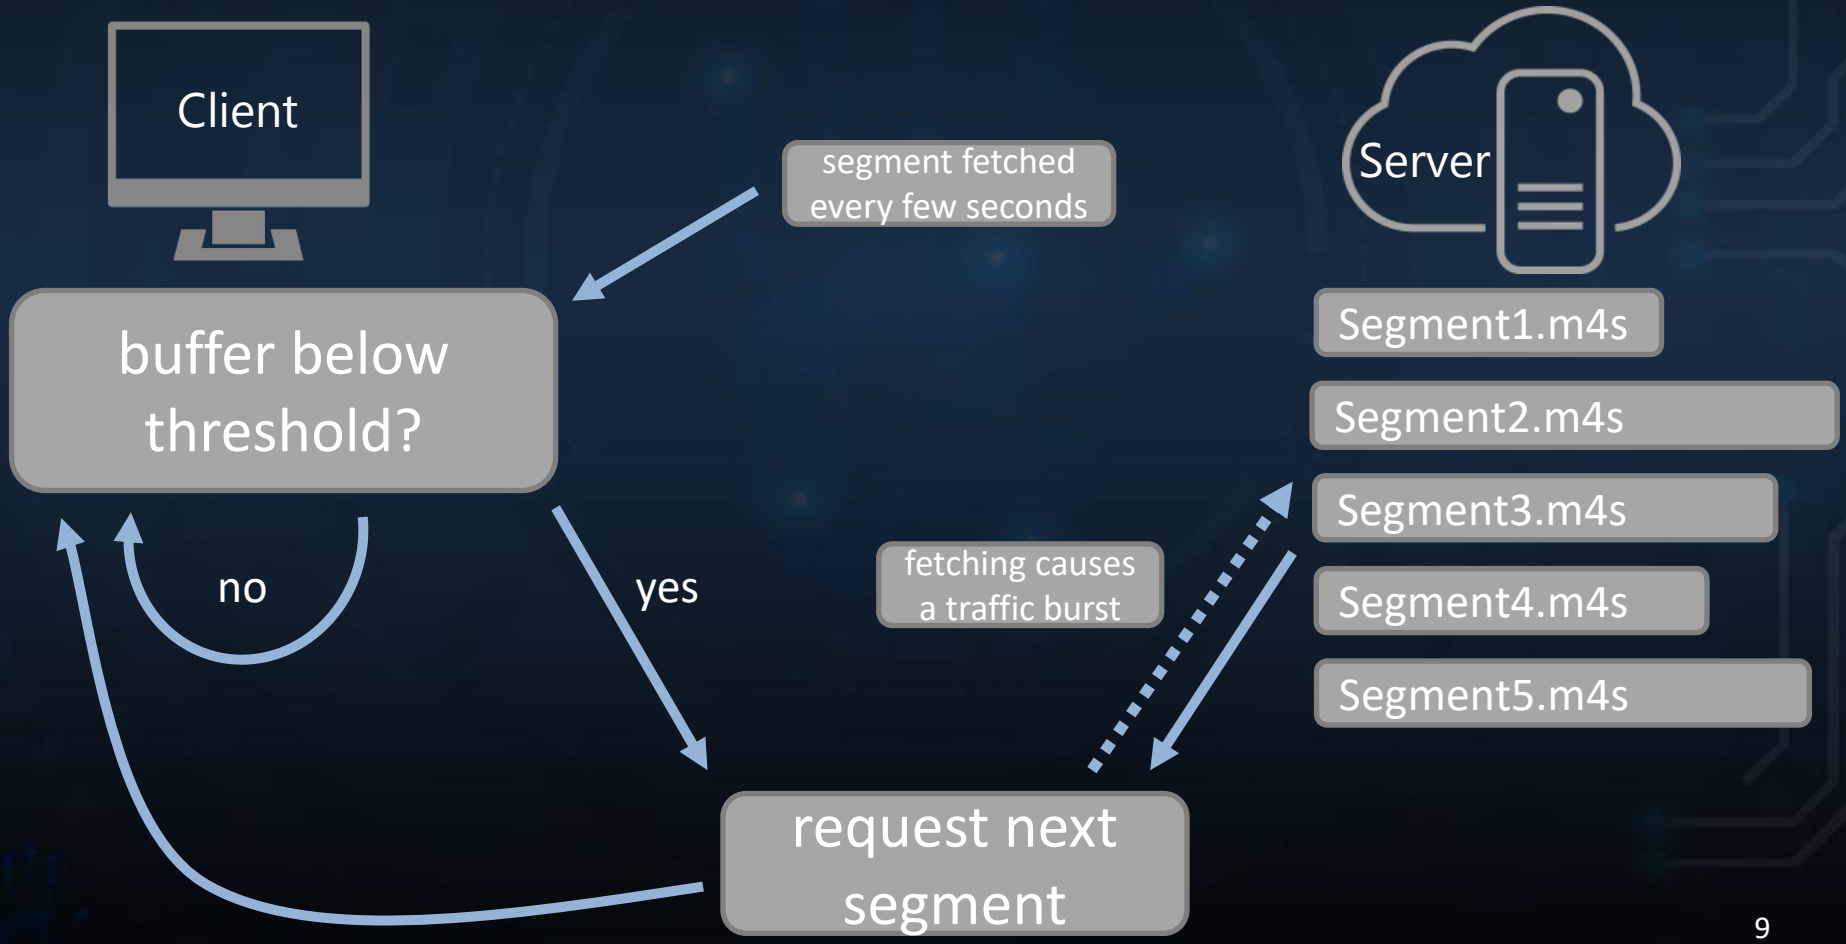

VBR Demo

Iguana vs. Snakes

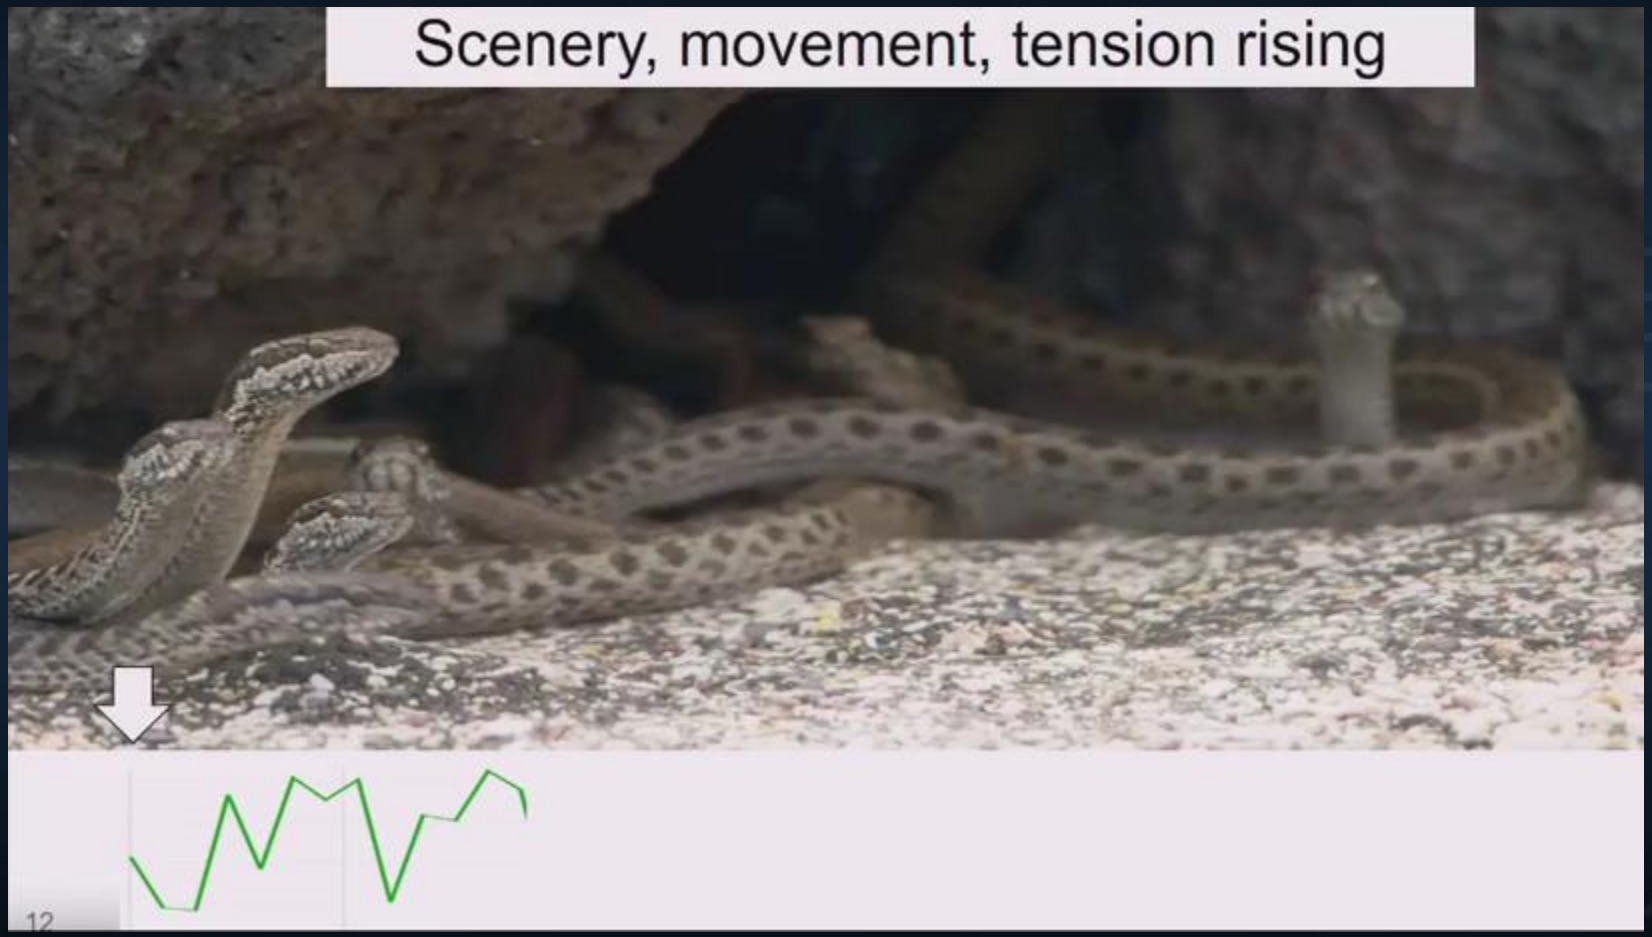

MPEG-DASH leak

Content

Burst sizes

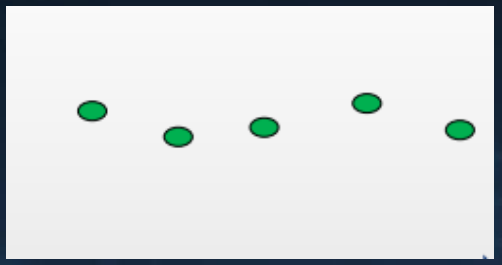

VBR pattern

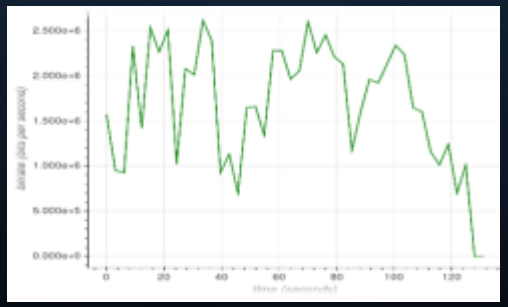

Segments

- Segment1.m4s
- Segment2.m4s
- Segment3.m4s

From a leak to a fingerprint

- Does the pattern of burst (segment) sizes uniquely characterize a title?
- Empirically for the 3500 downloaded from YouTube 20% of them have a uniquely identifying pattern
- Can we learn a title's identifying pattern?
- We can learn a title's identifying pattern because of the pattern consistency

~20% of YouTube titles have fingerprints

Attack overview

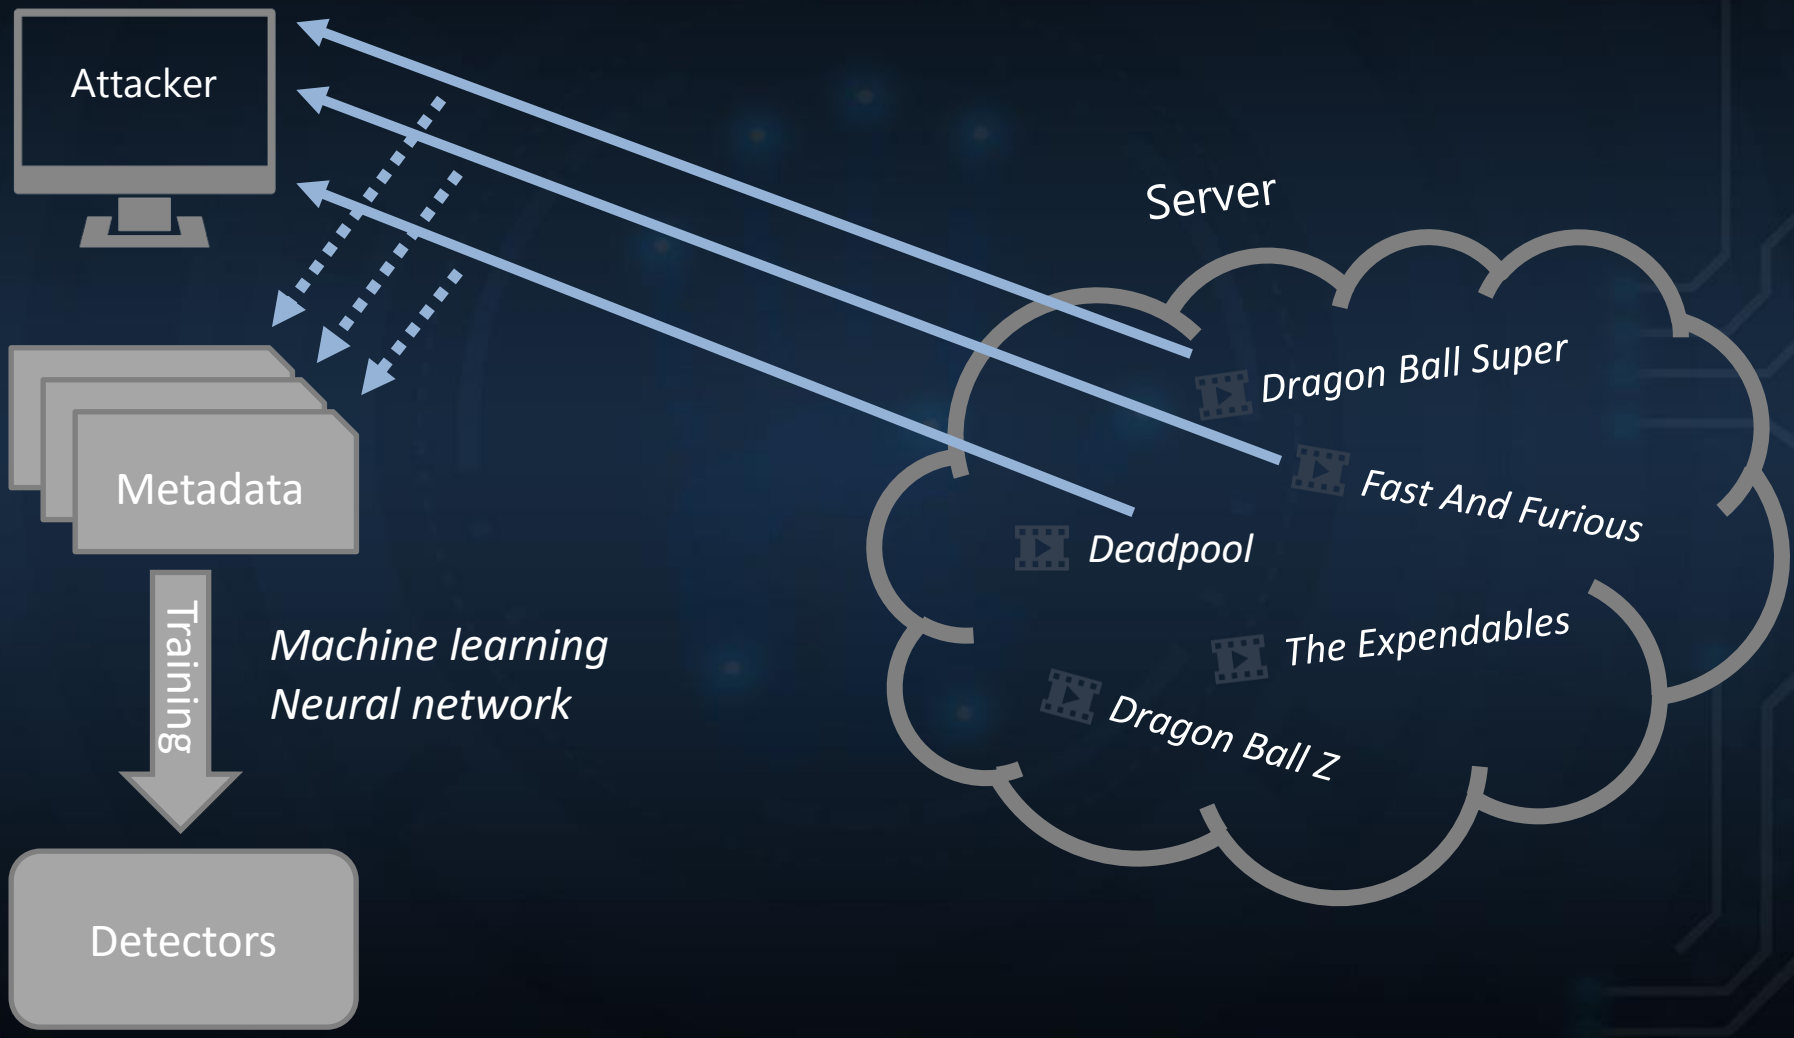

Attack overview

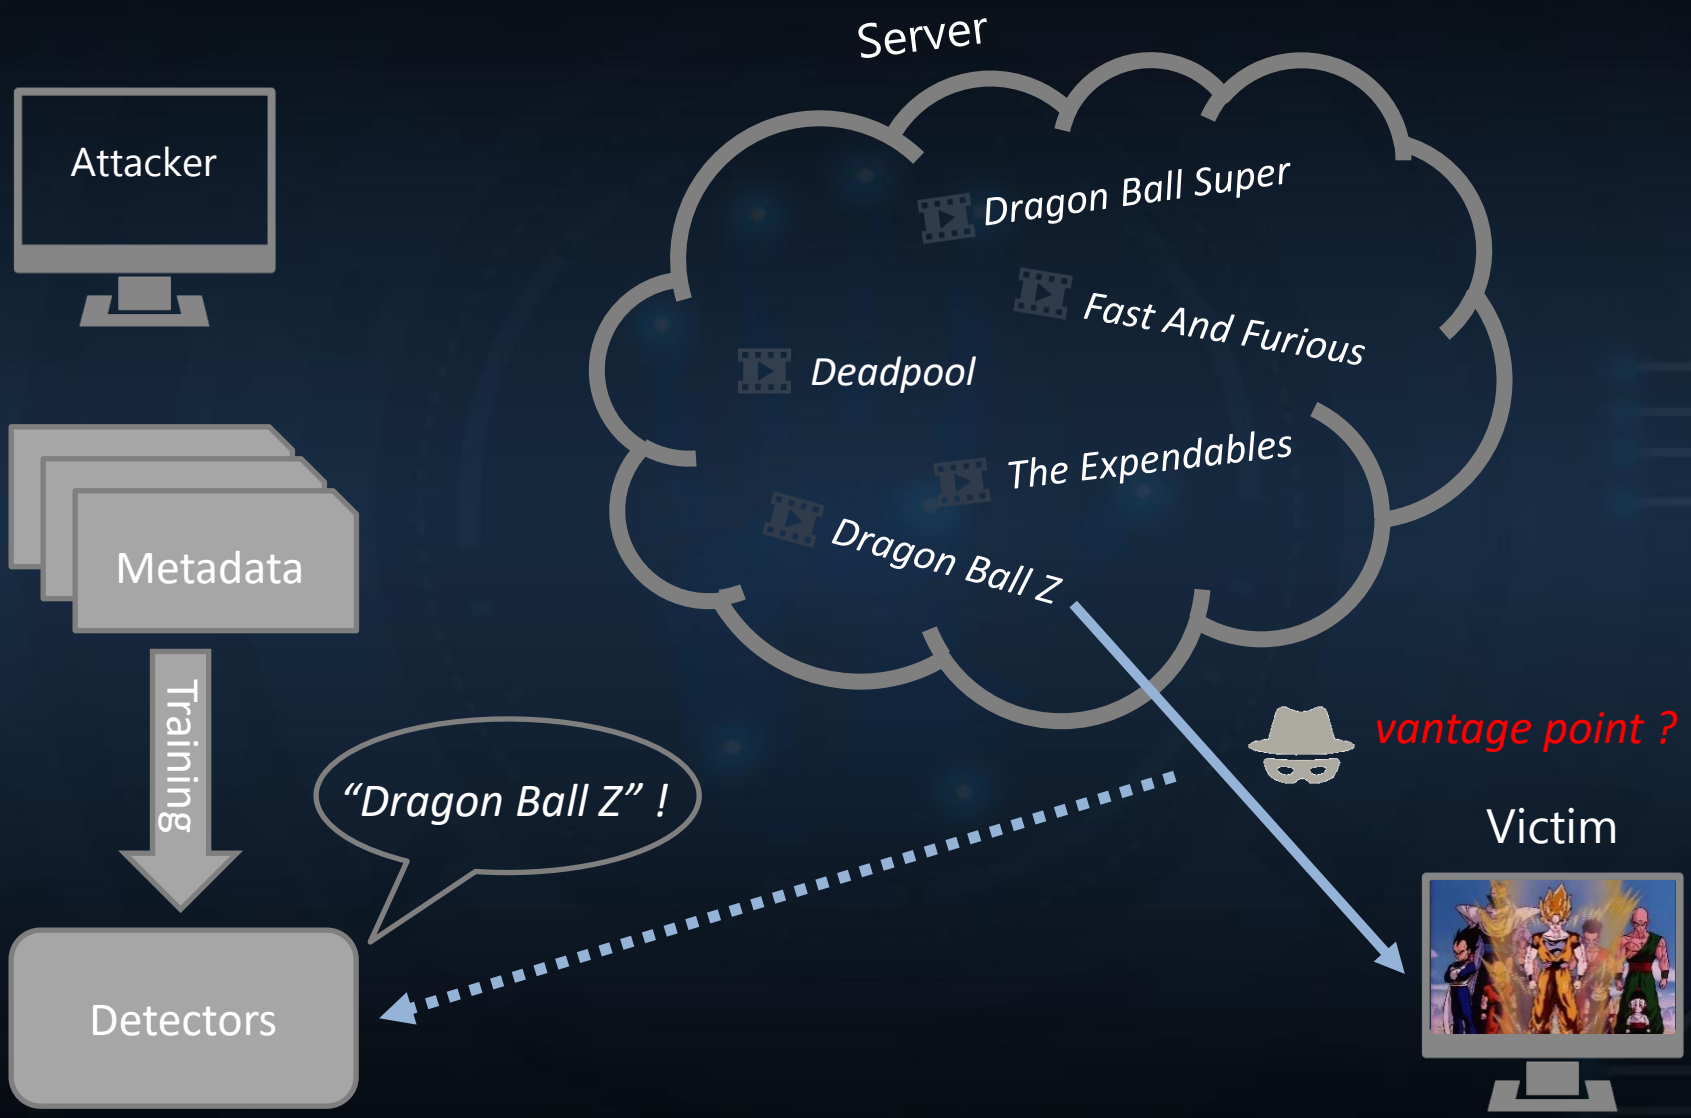

Deep Neural Networks

- Very good at learning high-level concepts that human can easily agree on but find it hard to formally express
- Can operate on noisy and coarse measurements
- Agnostic to protocol-specific attributes
- Can learn features other than burst patterns, e.g. arrival patterns of individual packets
- Can use multiple session representations, train on all at once

Deep Neural Networks

100 titles
98.5% accuracy

YouTube

18 titles + 3500 sessions
of different other titles
99.5% accuracy

amazon instant video

10 titles
92.5% accuracy

vimeo

10 titles
98.6% accuracy

Tuning for precision

YouTube

feature: total burst size

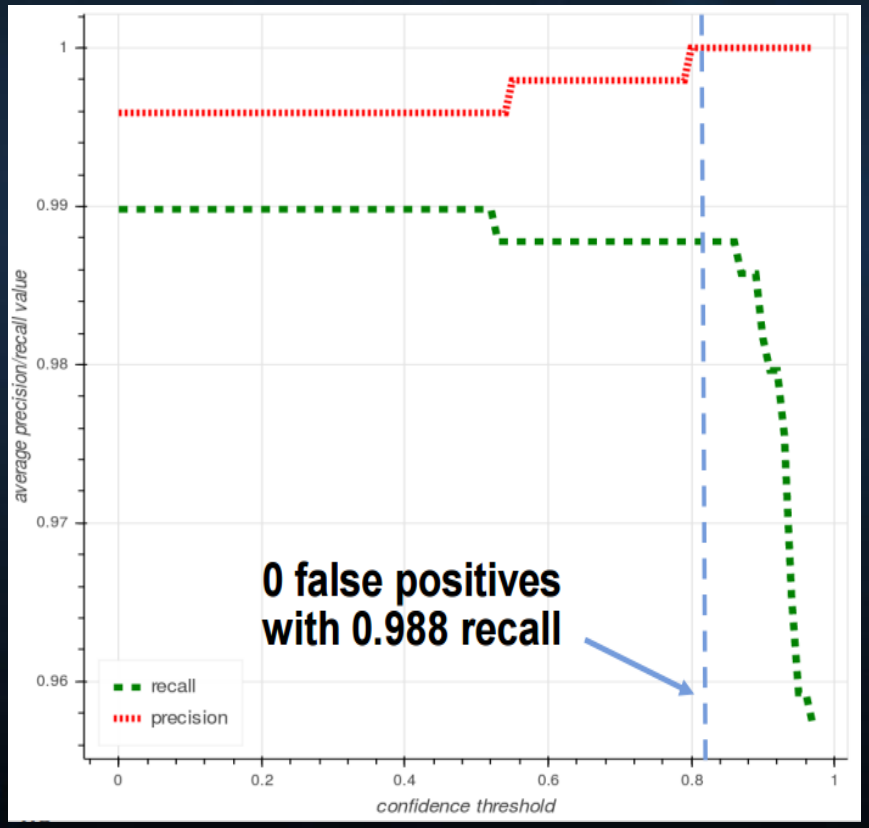

Netflix

feature: total burst size

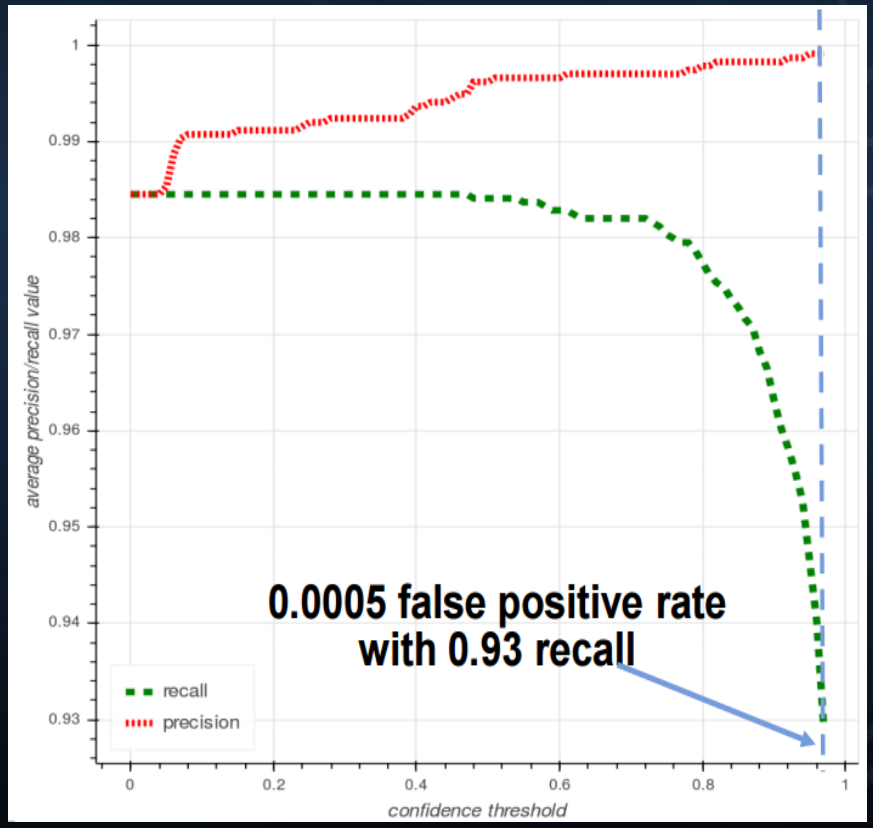

So what is the vantage point ?

Scenario I: on-path attack

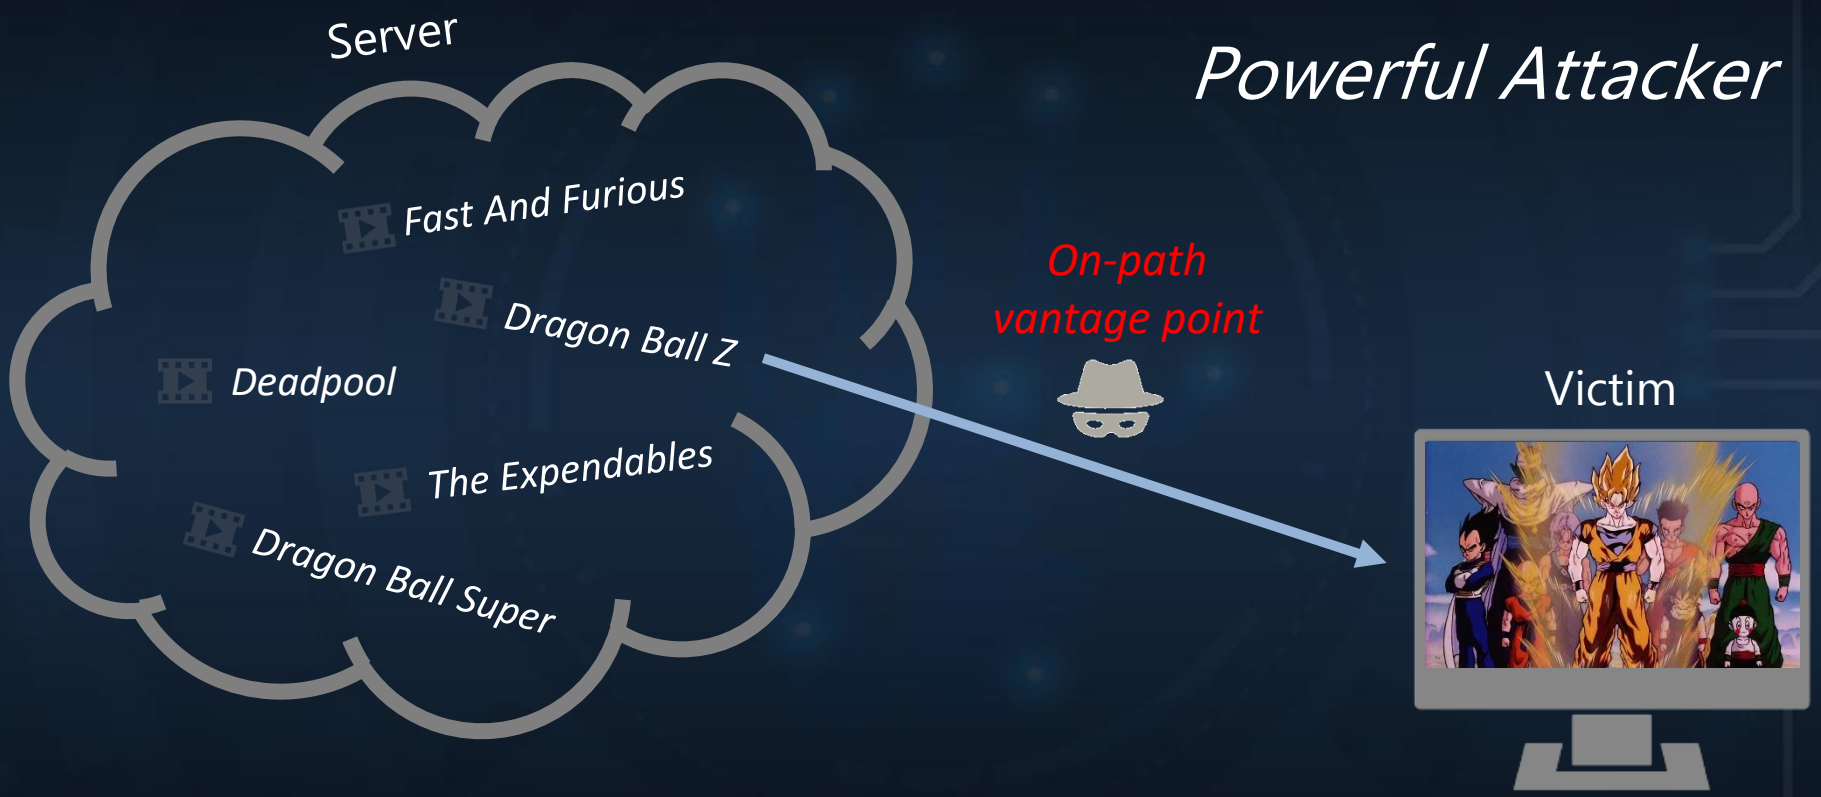

Wi-Fi access points, proxies, routers, enterprise or national network censors, ISPs

Scenario II: off-path attack

Server

Dragon Ball Super

Fast And Furious

Deadpool

The Expendables

Dragon Ball Z

Off-path vantage point

Victim network

Malicious AD

Checking Facebook feed while streaming "Dragon Ball Super"

Cross-device attack

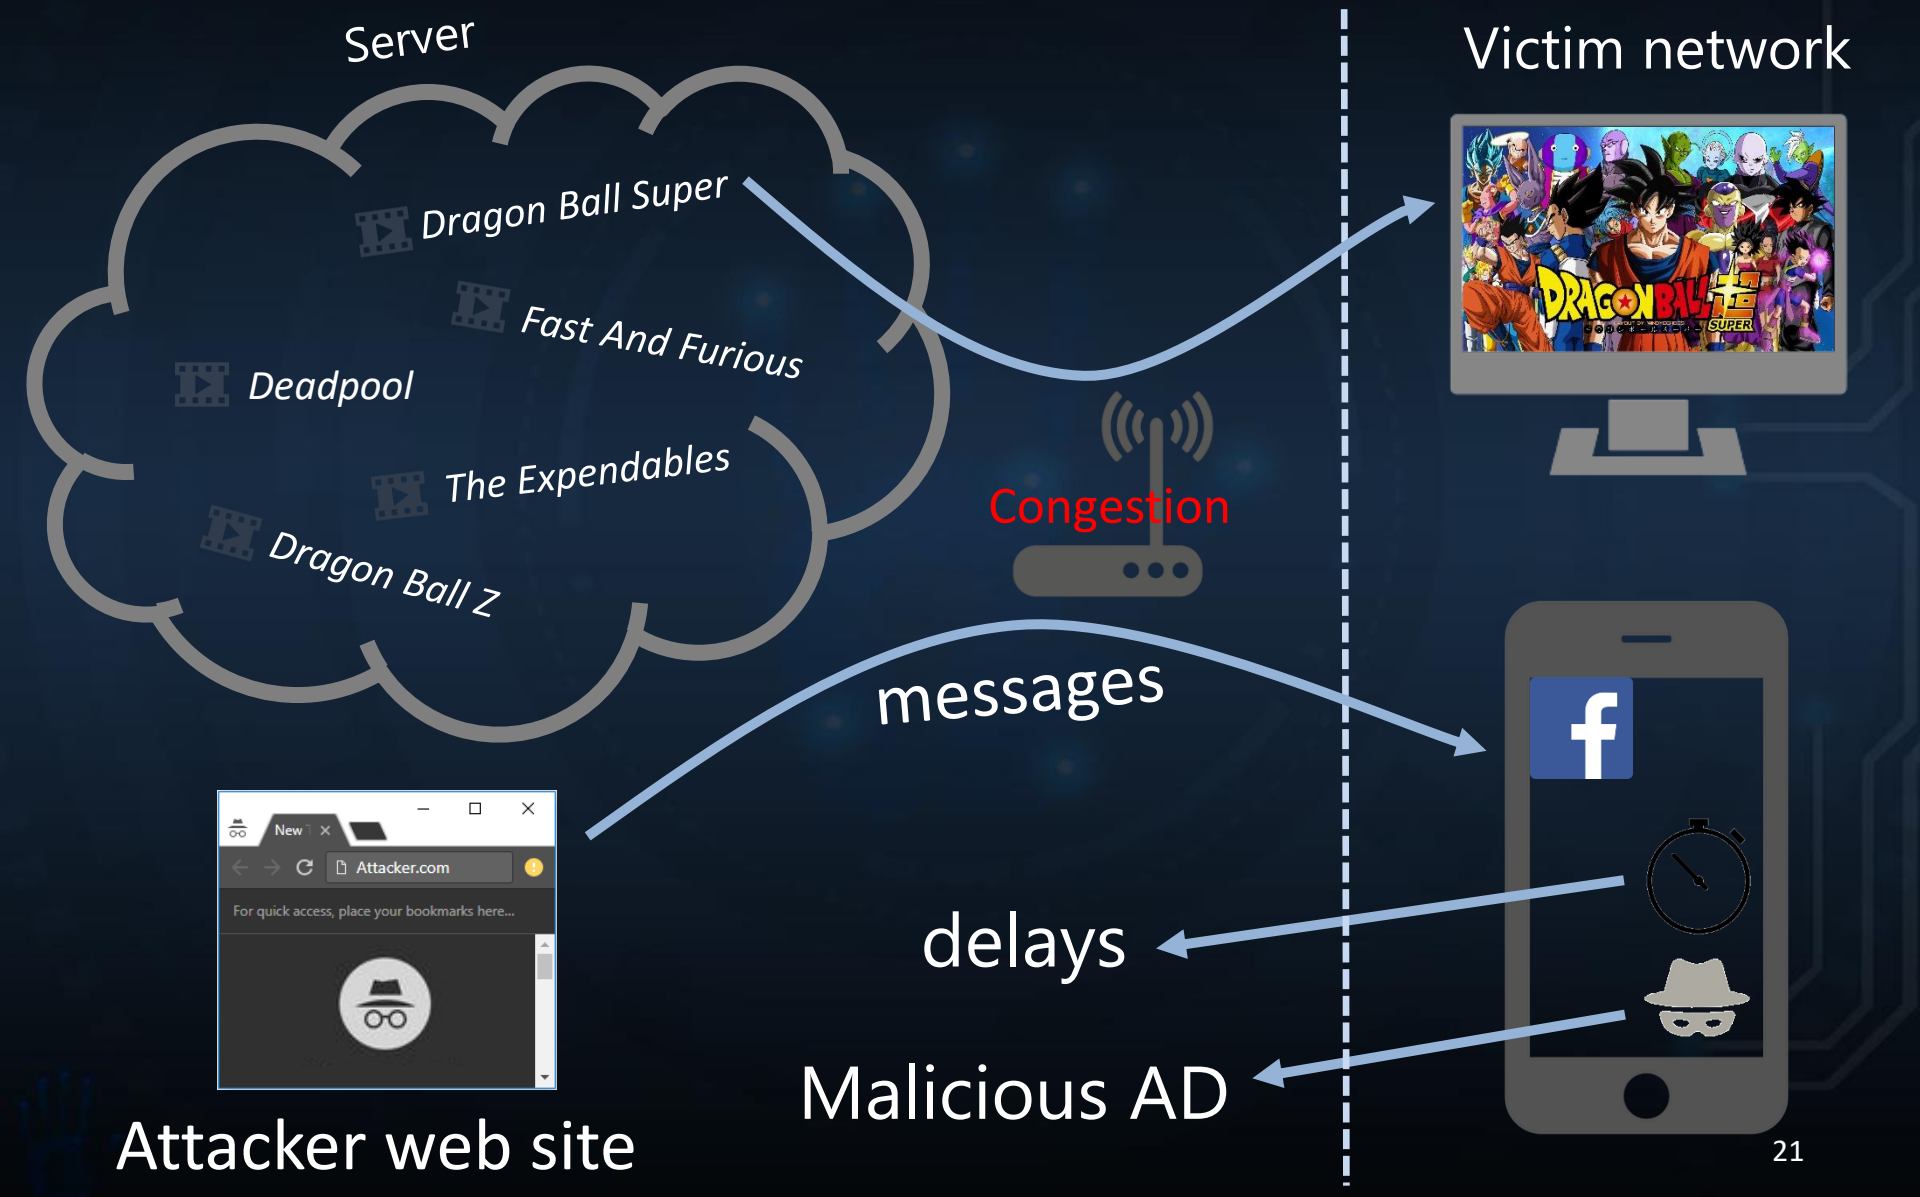

Delay-bursts

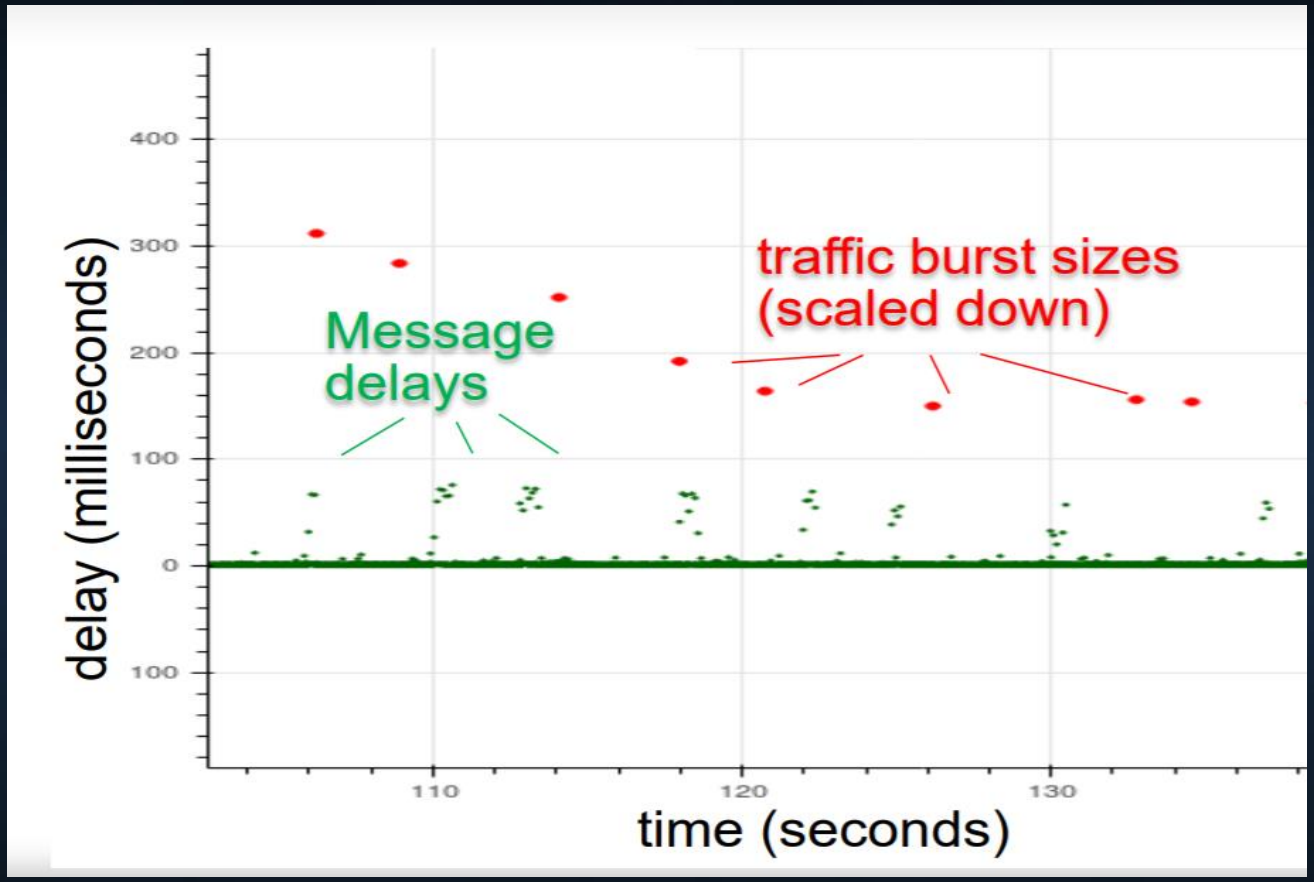

For each traffic burst, compute aggregate delay induced
Use resulting time-series as input to neural network

Delay-bursts VS. Traffic-bursts

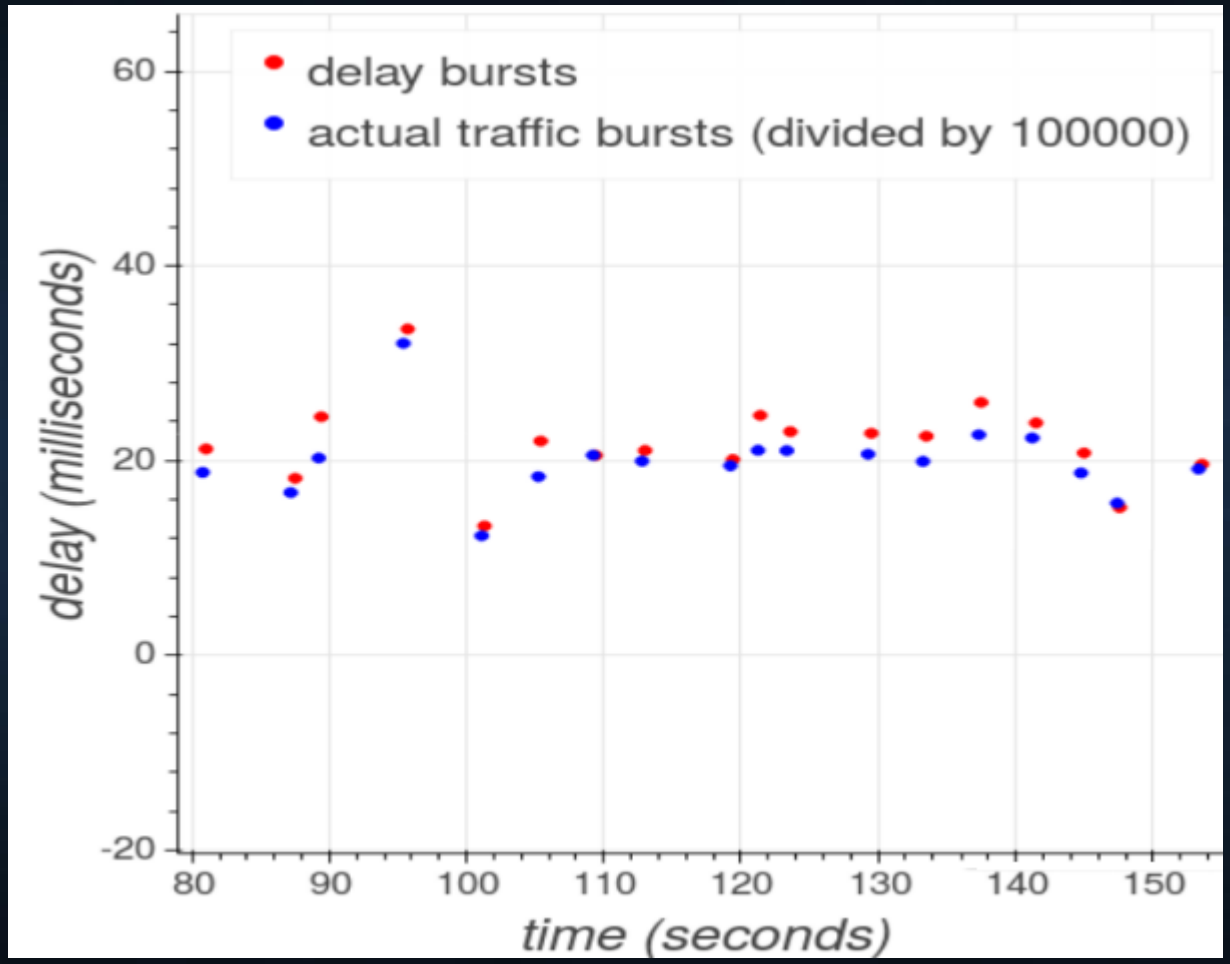

Delay-bursts time series: the delays induced by traffic bursts

Cross-site attack

Mitigating the MPEG-DASH leak

- Modern streaming traffic characteristics
  - Title bitrate pattern unique when sampled at few-seconds granularity
  - Fetching at segment granularity (every few seconds)
- Maximizes quality of experience (QoE), server load, and network bandwidth utilization
- However, information leakage is intrinsic…

Conclusions

- Leakage of information about video content via network traffic patterns is prevalent in modern streaming protocols and popular services
- Detectors are tuned for high accuracy and effective in an “open-world” setting
- It can be used by on-path adversaries such as ISPs to spy on their users
- It can be used by off-path adversary who merely serves a Web page to identify videos being streamed by the user

Thank You !

- Further information and the paper: https://beautyburst.github.io/

"Everything has a fingerprint,
and so do encrypted streams"

Any Questions?